Analysis of optimization and numerical approaches to solve the linear least square problem

Emanuele Cosenza*, Riccardo Massidda†

Department of Computer Science
University of Pisa

* e.cosenza3@studenti.unipi.it, † r.massidda@studenti.unipi.it

Abstract—**The linear least square problem can be tackled using a wide range of optimization or numerical methods. The L-BFGS method of the class of limited-memory quasi-Newton algorithms has been chosen for the former, whilst the thin QR factorization with Householder reflectors for the latter. Both these algorithms have been implemented from scratch using Python language, to finally experiment over their performances in terms of precision, stability and speed. The accordance of the implementations with the underlying theoretical models is also studied and discussed.**

Introduction

Given a dataset composed by a matrix $\hat{X} \in \mathbb{R}^{m \times n}$ with $m \geq n$ and a vector $y \in \mathbb{R}^m$, the solution of the linear least square (LLS) problem is the vector $w \in \mathbb{R}^n$ that fits best the data assuming a linear function between $\hat{X}$ and y. (Nocedal and Wright 2006, 50) This can be formalized as the following minimization problem:

$$w_* = \min_w \|\hat{X}w - y\|_2^2$$

The matrix $\hat{X}$ is actually composed in the following way:

$$\hat{X} = \begin{bmatrix} X^T \\ I \end{bmatrix}$$

Where $X \in \mathbb{R}^{n \times k}$ is a tall thin matrix, thus $m = k + n$. The LLS problem can be dealt both with iterative methods or with direct numerical methods. One algorithm has been chosen for each of these fields to finally discuss their experimental results.

L-BFGS

The Limited-memory BFGS, L-BFGS, is an iterative method of the quasi-Newton limited-memory class. This method is a variation of the BFGS method, with which it shares the update rule. At the $i + 1$-th iteration the point is updated as follows:

$$w_{i+1} = w_i - \alpha_i H_i \nabla f_i$$

The smaller memory requirements of this variation are due to the fact that the Hessian approximation H_i is stored implicitly, and built over a fixed number of vector pairs $\{s_j, y_j\}$ of the previous t iterations and an initial matrix H_i^0. Where

$$s_i = w_{i+1} - w_i, \quad y_i = \nabla f_{i+1} - \nabla f_i$$

$$V_i = I - \rho_i y_i s_i^T, \quad \rho_i = \frac{1}{y_k^T s_k}$$

so that H_i satisfies the following condition

$$
\begin{aligned}
H_i =\ & (V_{i-1}^T \ldots V_{i-t}^T) H_i^0 (V_{i-t} \ldots V_{i-1}) \\
& + \rho_{i-t}(V_{i-1}^T \ldots V_{i-t}^T + 1) s_{i-t} s_{i-m}^T (V_{i-t+1} \ldots V_{i-1}) \\
& + \rho_{i-t+1}(V_{i-1}^T \ldots V_{i-t}^T + 2) s_{i-t+1} s_{i-t+1}^T (V_{i-t+2} \ldots V_{i-1}) \\
& + \ldots \\
& + \rho_{i-1} s_{i-1} s_{i-1}^T
\end{aligned}
$$

Different strategies to initialize the H_i^0 matrix are proposed in the literature, and so they will be tested experimentally. Finally, the step size α_i is found by performing an inexact line search based on the Armijo-Wolfe conditions.

Thin QR factorization

For the numerical counterpart, the thin QR factorization with Householder reflectors has been implemented as described in (Trefethen and Bau 1997).

By using the Householder QR factorization, the matrix $R \in \mathbb{R}^{m \times n}$ is constructed in place of $\hat{X}$ and the n reflection vectors $v_1, \ldots, v_n$ are stored. The reduced matrix $\hat{R} \in \mathbb{R}^{n \times n}$ is trivially obtainable by slicing as in $\hat{R} = R_{1:n,1:n}$, noting that constructing directly the reduced matrix would yield no significant advantage since we are already using $O(mn)$ space to store the reflection vectors.

By using the Householder vectors it is also possible to implicitly compute $\hat{Q}^T y$ to finally obtain w_* by back substitution over the upper-triangular system $\hat{R}w = \hat{Q}^T y$.

Algorithmic analysis

Convergence of L-BFGS

Liu and Nocedal (1989) define three necessary assumptions to prove that the L-BFGS algorithm globally converges and that there exists a constant $0 \leq r < 1$ such that

$$f(w_i) - f(w_*) \leq r^i(f(w_0) - f(w_*))$$

so that the sequence $\{w_i\}$ converges R-linearly.

Firstly the objective function f should be twice continuously differentiable. Given the formulation of the least squares problem this is immediately true, the gradient and the Hessian of the objective function are definable as in:

$$\nabla f(w) = \hat{X}^T(\hat{X}w - y)$$

$$\nabla^2 f(w) = \hat{X}^T \hat{X}$$

Moreover the Hessian can be proven to be positive definite, as can be easily seen by rearranging it in the following way:

$$\begin{aligned} \nabla^2 f(w) &= \hat{X}^T \hat{X} \\ &= \begin{vmatrix} XI \end{vmatrix} \begin{vmatrix} X^T \\ I \end{vmatrix} \\ &= XX^T + I \end{aligned}$$

The matrix XX^T is positive semi-definite, since $\forall z : z^T XX^T z = \|X^T z\| \geq 0$, therefore all the eigenvalues of the matrix are non-negative. Furthermore, according to the spectral theorem, since XX^T is symmetric, there exists U orthogonal matrix and D diagonal containing the eigenvalues of XX^T.

$$\begin{aligned} \nabla^2 f(x) &= XX^T + I \\ &= UDU^T + I \\ &= UDU^T + UIU^T \\ &= U(D+I)U^T \end{aligned}$$

The eigenvalues of the Hessian are contained in $D + I$ and all of them are positive, therefore $\nabla^2 f(w)$ is positive definite.

Being the Hessian positive definite, the objective function f is a convex function. This comes in handy for the second assumption requiring the sublevel set $D = \{w \in \mathbb{R}^n | f(w) \leq f(w_0)\}$ to be convex. It can be easily proved that if a function is convex all of its sublevel sets are convex sets.

$$\forall x, y \in D, \lambda \in [0, 1]$$
$$\text{f convex}$$
$$\begin{aligned} &\implies f(\lambda x + (1-\lambda)y) \\ &\leq \lambda f(x) + (1-\lambda)f(y) \\ &\leq \lambda f(w_0) + (1-\lambda)f(w_0) \\ &= f(w_0) \\ &\implies \lambda x + (1-\lambda)y \in D \end{aligned}$$

The third and last assumption requires the existence of two positive constants M_1 and M_2 such that $\forall z \in \mathbb{R}^n, w \in D$:

$$M_1\|z\|^2 \leq z^T \nabla^2 f(w)z \leq M_2\|z\|^2$$

or equivalently

$$M_1 I \preceq \nabla^2 f(w) \preceq M_2 I$$

Since $\nabla^2 f(w)$ is positive definite the previous condition is true for $M_1 = \lambda_{min}$ and $M_2 = \lambda_{max}$, where $\lambda_{min} > 0$.

In the convergence proof the M_2 constant is used to upper bound the trace of the next Hessian substitute H_{i+1}, implying an upper bound for the largest eigenvalue in the sequence of Hessian substitutes.

$$tr(H_{i+1}) \leq tr(H_i^0) + tM_2 \leq M_3$$

On the other hand the M_1 constant is used, to lower bound the determinant of H_{i+1}, implying a lower bound for the smallest eigenvalue in the sequence of Hessian substitutes.

$$det(H_{i+1}) \geq det(H_i^0) + (\frac{M_1}{M_3})^t \geq M_4$$

These two assertions are used to prove the existence of constant $\delta > 0$ such that

$$\forall i : \cos \theta_i = \frac{s_i^T H_i s_i}{\|s_i\|\|H_i s_i\|} \geq \delta$$

where θ_i is the angle between the chosen direction and $-\nabla f(w_i)$. If the constant M_1 was to be equal to zero, it would not be enough to prove the existence of $\delta > 0$ for each step, possibly having directions orthogonal to steepest one. As already pointed out, given that the Hessian is positive definite, its eigenvalues and consequently M_1 are positive.

Other then the three discussed assumptions, the theorem requires for the sequence of initializers $\{\|H_i^0\|\}$ to be bounded. This obviously depends on the initialization technique used to generate H_i^0, various techniques are suggested in the literature such as $H_k^0 = \gamma_k I$ or $H_k^0 = \gamma_k H_0$ where

$$\gamma_k = \frac{s_{k-1}^T y_{k-1}}{\|y_{k-1}\|}$$

Other initialization techniques may possibly be tested and evaluated experimentally.

Armijo-Wolfe inexact line search

The convergence proof requires the algorithm to perform a line search respectful of the Armijo-Wolfe conditions, the solution described in Al-Baali and Fletcher (1986) is therefore adapted and implemented.

The algorithm performs an inexact line search that is ensured to converge under the assumption that $\sigma > \rho$ where $\rho \in (0, \frac{1}{2}), \sigma \in (0, 1)$, respectively the constant for the Armijo condition and for the Wolfe one. By defining the function ϕ, used to evaluate the value of f at a certain step-size α, the conditions can be defined as follows.

$$\phi(\alpha) = f(w_i + \alpha d_i)$$

$$\phi(\alpha) \leq \phi(0) + \alpha\rho\phi'(0) \tag{A}$$

$$\phi'(\alpha) \geq \sigma\phi'(0) \tag{W}$$

The algorithm requires a lower bound $\bar{f}$ on $\phi(\alpha)$ for $\alpha \geq 0$. More precisely, it assumes that the user is prepared to accept any value of α for which $\phi(\alpha) \leq \bar{f}$ where $\bar{f} < \phi(0)$. For the linear least-squares problem an obvious lower bound is $\bar{f} = 0$.

The algorithm performs an inexact line search by looking for a candidate point α_i at the i-th iteration in the interval (a_i, b_i), stopping if such candidate reaches the lower bound or if it satisfies both (A) and (W).


Fig. 1. Graphical depiction of the μ point.

The Armijo condition describes a line, called ρ-line, in the plot $(\alpha, \phi(\alpha))$ that can be useful to bound the starting interval. In fact the initial search interval can be reduced from $(0, \infty)$ to $(0, \mu)$ where

$$\mu = \frac{\bar{f} - \phi(0)}{\rho\phi'(0)}$$

It is immediate that $\forall \alpha > \mu$ either (A) can't be satisfied, or the point lies under the lower bound $\bar{f}$ (figure 1).

To proceed with the discussion over the shrinking procedure the function T is defined as in

$$T(a, b) = [a + \tau_1(b - a), b - \tau_2(b - a)]$$

where $0 < \tau_1 \leq \tau_2 \leq \frac{1}{2}$.

If the candidate doesn't satisfy (A) or if the left extreme a_i constitutes a better point, the next candidate is chosen in the interval $T(a_i, \alpha_i)$. Otherwise if the candidate doesn't satisfies (W) the next candidate is chosen in $T(\alpha_i, b_i)$. In both cases the a_{i+1} and b_{i+1} are updated with the extremes returned by the T function.

The candidate step-size may be randomly chosen between all the points in the interval defined by the T function, this approach will be experimentally tested against quadratic interpolation.

It should be noted that the Al-Baali and Fletcher (1986) paper defines the function T in a slightly different way, together with another function E used to specifically define the interval when (W) is not satisfied. The simplification hereby described is due to the fact that in our implementation it is ensured that $\forall i : a_i \leq \alpha_i \leq b_i \land b_i \neq \infty$, moreover this does not interfere with the convergence proof.

As suggested by Liu and Nocedal (1989) the unitary step length should always be tried first, so the first candidate should be $\alpha_0 = 1$. Other suggestions known in literature about the initialization of the remaining hyper-parameters are presented in the experimental setup section to be eventually evaluated.

Analysis of standard and modified thin QR

Ignoring constants, we know from theory that the standard QR factorization algorithm applied on the matrix $\hat{X}$ yields a time complexity of $O(mn^2)$. Actually, given that we are generally dealing with a very tall and thin matrix X, the resulting $\hat{X}$ is going to be squarish ($m \approx n$). This means that we can consider the complexity to be cubic in n.

From now on, we show a way to bring down the time complexity of the algorithm from $O(mn^2)$ to $O(kn^2)$, with $k = m - n$. The resulting modified QR factorization algorithm will become useful when k is much smaller than m, as in our case.

In the standard algorithm, at each step we focus on a single column of the input matrix, constructing a householder vector to zero out all the entries below the diagonal. Following the geometric reasoning in (Trefethen and Bau 1997), this brings the algorithm to depend on m. While this cannot be avoided in general, in our particular case we can be a little bit smarter.

Since the block matrix $\hat{X}$ contains the identity as its lower block, at each step of the algorithm we can just focus on

zeroing out the $k = m - n$ entries below the diagonal up to the 1s of the identity. In the modified algorithm then, to obtain R, the matrix $\hat{X}$ is multiplied on the left side by a sequence of matrices $L_i \in \mathbb{R}^{m \times m}$ of the form ($i = 1, \ldots, n$):

$$L_i = \begin{bmatrix} I_{i-1} & 0 & 0 \\ 0 & H_i & 0 \\ 0 & 0 & I_{n-i} \end{bmatrix}$$

where $H_i \in \mathbb{R}^{(k+1) \times (k+1)}$ are all Householder reflectors that zero out the k entries in the i-th column of the matrix which is being multiplied by L_i. The resulting Householder vectors will all have constant dimension $k + 1$ and storing them will require $\theta(kn)$ space instead of $O(mn)$.

To derive the time complexity of this phase we can reason as follows. The right side matrix can be divided in three blocks as in $\begin{bmatrix} A \\ B \\ C \end{bmatrix}$. When we multiply this matrix by L_i the only relevant operation is the matrix multiplication $H_i B$, which costs $O(kn)$. Since the total number of multiplications is n, the total complexity is $O(kn^2)$.

Since each L_i is orthogonal and symmetric it is then possible to reconstruct Q and the reduced $\hat{Q}$ in the following way:

$$Q = L_1 L_2 \ldots L_n$$
$$\hat{Q} = L_1 L_2 \ldots L_n \begin{bmatrix} I_n \\ 0 \end{bmatrix}$$

Applying again the reasoning above, the time complexity of these reconstructions is $O(kn^2)$. It follows that the overall time complexity of the modified QR factorization is $O(kn^2)$.

The least squares problem is then solved through back substitution over the upper-triangular system $\hat{R}w = \hat{Q}^T y$. Since the costs for the $\hat{Q}^T y$ product and the back substitution are dominated by the factorization cost, the overall time complexity for solving the least squares problem through QR factorization is $O(mn^2)$ when using the standard algorithm and $O(kn^2)$ when using the modified one.

Stability and accuracy of the QR algorithm

As stated in (Trefethen and Bau 1997, 140), the algorithm obtained by combining the standard QR algorithm, the $Q^T y$ product and back substitution is backward stable in the context of least squares problems.

We claim that the QR factorization step remains backward stable if we consider the modified version described in this report. Without going into details with an extended proof, this can be explained by saying that at each step of the algorithm we apply a transformation L_i doing a smaller number of operations than those of the standard algorithm. Then, since we know that each step of the standard QR factorization is backward stable, this must be true also in the modified version of the algorithm.

Considering that both versions of the QR algorithm are backward stable, the accuracy of the algorithms will depend mostly on the conditioning of the least squares problem at hand. In fact, following from the definition of backward stability, the algorithms will both find exact solutions to slightly perturbed problems, with perturbations of the order of machine precision. This implies that if the conditioning of the problem is high the real solutions to the perturbed problems are inevitably going to be inaccurate.

If w_* is the exact solution to the least squares problem and $\tilde{w}_*$ is the solution found with one of the QR based algorithms outlined above, the accuracy of the algorithms will therefore follow the general upper bounds of relative errors found in (Trefethen and Bau 1997, 131): one relative to perturbations of the matrix $\hat{X}$,

$$\frac{\|\tilde{w}_* - w_*\|}{\|w_*\|} \leq (\kappa(\hat{X}) + \kappa(\hat{X})^2 \tan \theta) \frac{\|\delta \hat{X}\|}{\|\hat{X}\|}$$

and one relative to perturbations of the vector y,

$$\frac{\|\tilde{w}_* - w_*\|}{\|w_*\|} \leq \left(\frac{\kappa(\hat{X})}{\cos \theta} \right) \frac{\|\delta y\|}{\|y\|}$$

where θ is the angle such that $\cos \theta = \frac{\|\hat{X} w_*\|}{\|y\|}$.

From these upper bounds we can expect that the algorithm will be more accurate when the angle θ is near 0 and less accurate when it is near $\frac{\pi}{2}$, reminding that in our context the value of θ will depend on the value of the random vector y.

<div align="center">

INPUT DATA

IMPLEMENTATION DETAILS

EXPERIMENTAL RESULTS

CONCLUSIONS

BIBLIOGRAPHY

</div>

Al-Baali, M., and R. Fletcher. 1986. "An Efficient Line Search for Nonlinear Least Squares." *Journal of Optimization Theory and Applications* 48 (3): 359–77. https://doi.org/10.1007/BF00940566.

Liu, Dong C., and Jorge Nocedal. 1989. "On the Limited Memory BFGS Method for Large Scale Optimization." *Mathematical Programming* 45 (1-3): 503–28. https://doi.org/10.1007/BF01589116.

Nocedal, Jorge, and Stephen J. Wright. 2006. *Numerical Optimization.* 2nd ed. Springer Series in Operations Research. New York: Springer.

Trefethen, Lloyd N., and David Bau. 1997. *Numerical Linear Algebra.* Philadelphia: Society for Industrial; Applied Mathematics.